THE WHOLE PICTURE

An agentic platform, with its gaps visible

NEW: callers verified + evaluation gates deploys + agents propose fixes

03 AGENT INFRASTRUCTURE

Harness

- Context — BUILT
- Capabilities — BUILT
- Execution — BUILT / ALPHA
- Evaluation — BUILT (THIN)

Governance

- Non-human identity — BUILT
- Security controls — BUILT
- Agent observability — BUILT

⚠ Evaluation: one agent, eight questions

⚠ Chat UI is one shared identity

02 PATH DEFINITIONS

PARTIAL

- Deterministic
- Probabilistic
- Hybrid

Alert dispatch built; no path registry

01 TOOLING

BUILT

- Developer portal
- Delivery + resources
- Secrets + identity
- Edge + admission
- Data + telemetry

EVERYTHING AS CODE

BUILT

Reviewed files + shared checks

Eight paths built. A human still holds the merge and the approval.

THREE WAYS WORK MOVES

01 DETERMINISTIC • BUILT

Declared state is reconciled.

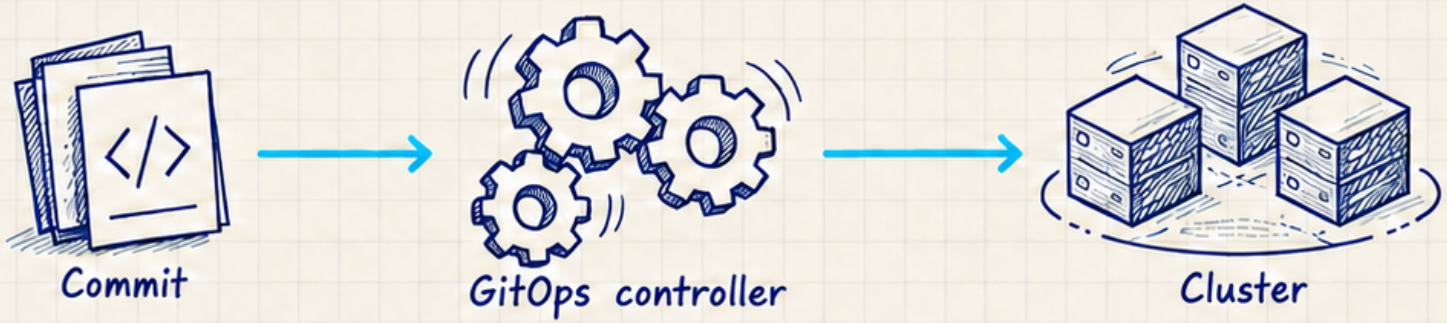

02 PROBABILISTIC • BUILT

Judgment helps route review.

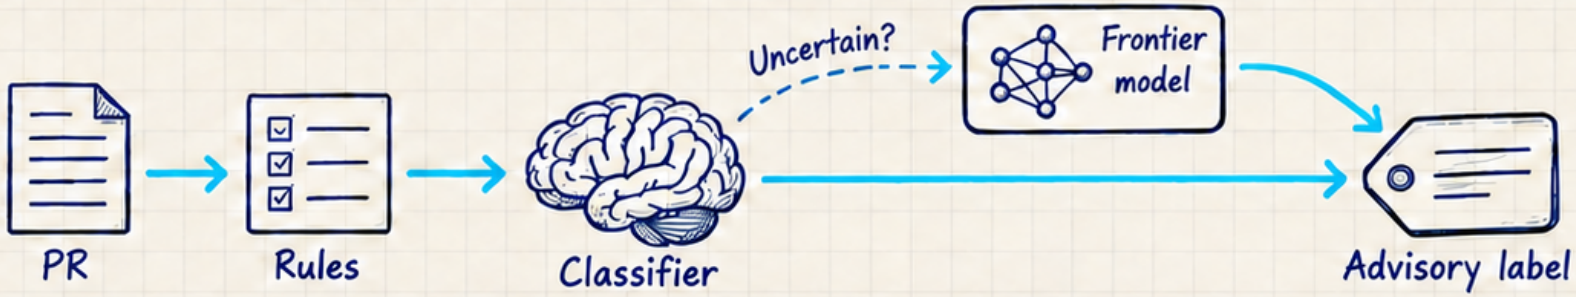

03 HYBRID • BUILT

Generation proposes. Checks constrain.

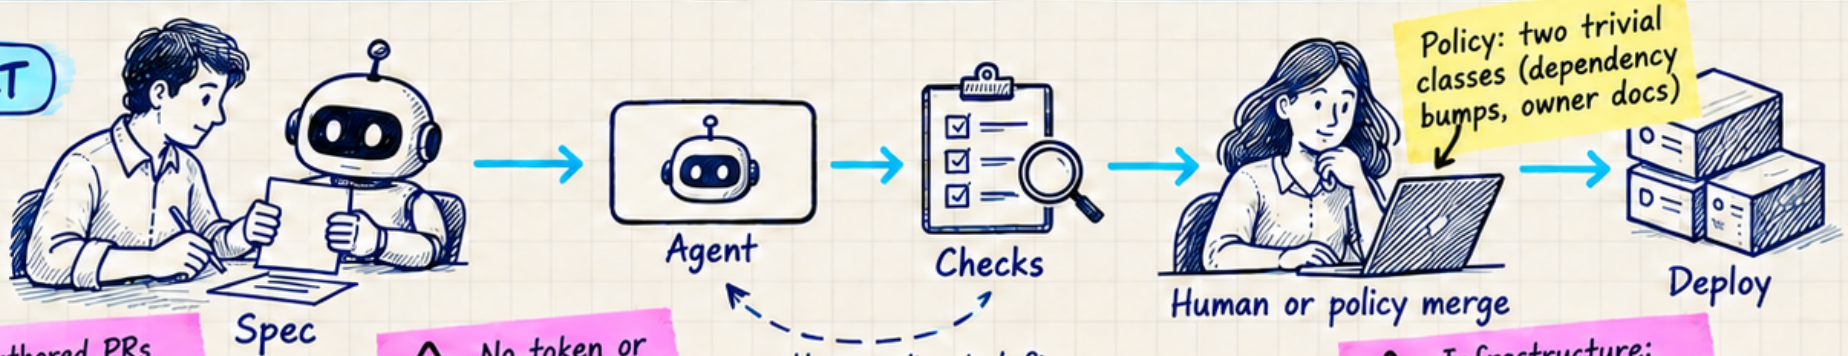

Agent-authored PRs always wait for a human.

⚠ No token or retry budget.

⚠ Infrastructure: apply before merge.

Checks constrain. Humans or policy promote.

GOVERNANCE SURROUNDS THE AGENT

IDENTITY • BUILT

- Scoped non-human principals
- Every caller verified at the gateway

Who may act?

What may happen?

SECURITY • BUILT

- Tool classes + approval gates
- Admission + CI enforcement

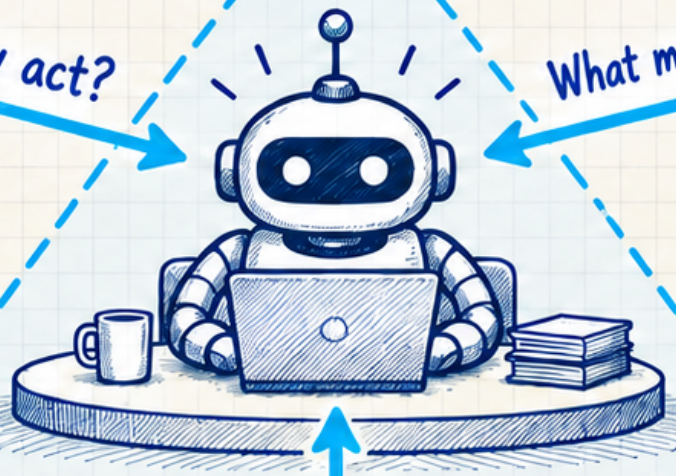

GAP: the chat UI is one identity

Agent-proposed restarts: the approver is recorded

What actually happened?

OBSERVABILITY • BUILT

- Calls + approvals + traces
- Daily audit + risk triage

Argument redaction is best-effort.

Declared policy must meet runtime evidence.

EIGHT PATHS, ONE HONEST SCORECARD

01 DISPATCH

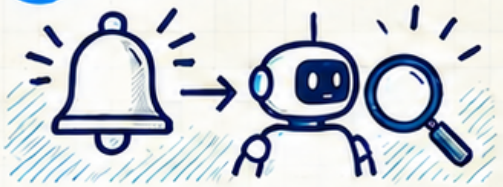

HYBRID | BUILT

Alerts → investigation, PR, or a held restart

02 RETRIEVE CONTEXT

PROBABILISTIC | BUILT

Docs, memory, contracts

03 IMPLEMENT

PROBABILISTIC | BUILT

Agent opens PRs under its own identity; a human merges

04 VALIDATE

HYBRID | BUILT

Checks + drills; fixes when asked

05 PROMOTE

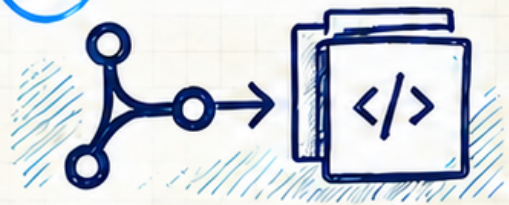

HYBRID | BUILT

Policy for two trivial classes; human for the rest

06 DEPLOY

DETERMINISTIC | BUILT

GitOps reconciliation

07 OBSERVE

HYBRID | BUILT

Alerts trigger findings in chat

08 REMEDIATE

PROBABILISTIC | BUILT

Limit too low → PR. Leak → a restart a human approves

8 BUILT • 0 PARTIAL • 0 NOT YET

Every path works.
The merge and the approval stay human.

How much autonomy exists?

01 Human in the loop

BUILT

Assistance on request.

02 Human as validator

BUILT

Agents write; humans or policy merge.

DAILY OPERATING MODE

03 Human orchestrates

PARTIAL

Desk subagents + alert dispatch.

A proposed restart is approved by a named person.

No agent-to-agent platform orchestration.

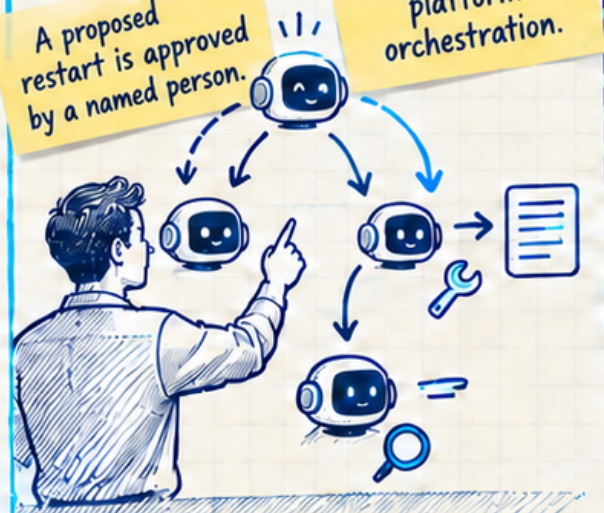

04 Human out of the loop

NOT YET

GENERAL CASE

Two declared exceptions: dependency bumps and owner docs merge by policy.

Two trivial classes only

Maturity belongs to each path, not to a blanket platform claim.

ONE RECORDED CHANGE

CLI approval flow • 12 OCT 2026

BEFORE

Human intent + agent implementation

9 files • +498 / −31 • 49 tests

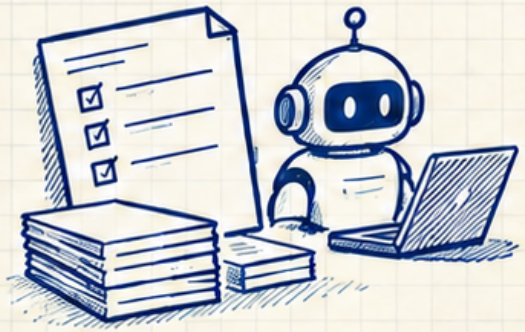

PREPARATION

23 MINUTES: PR OPEN → MERGE

15:59Z • PR opened

16:22Z • Human merge

Review classifier ran

27 check contexts • 4 skipped

Admission test timed out → human rerun passed

Admin override

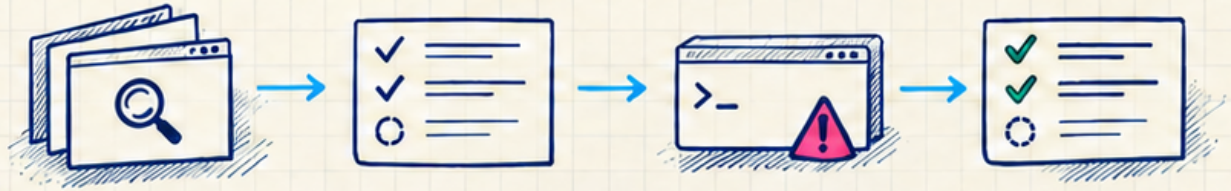

MEASURED WINDOW

AFTER

Reject: no change, no retry

Explicit request again → approve

Docs + memory updated

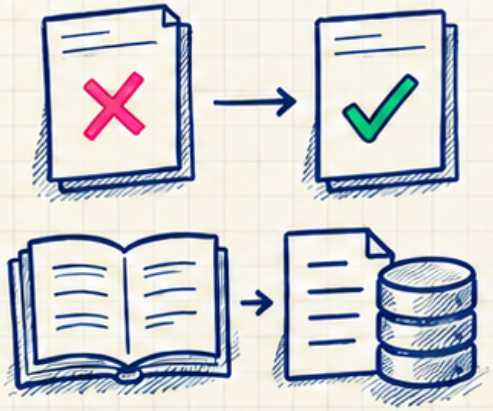

ACCEPTANCE

23 minutes excludes implementation and live acceptance.

A real delivery record, with human decisions still visible.

THE AUDIT THAT CHANGED THE CONTROLS

Historical finding • July 2026

GATED TOOLS ACTUALLY INVOKED

14
shell calls
09 JUL 2026

1
delete call
19 APR 2026

0
approval requests
Across the audited agent history

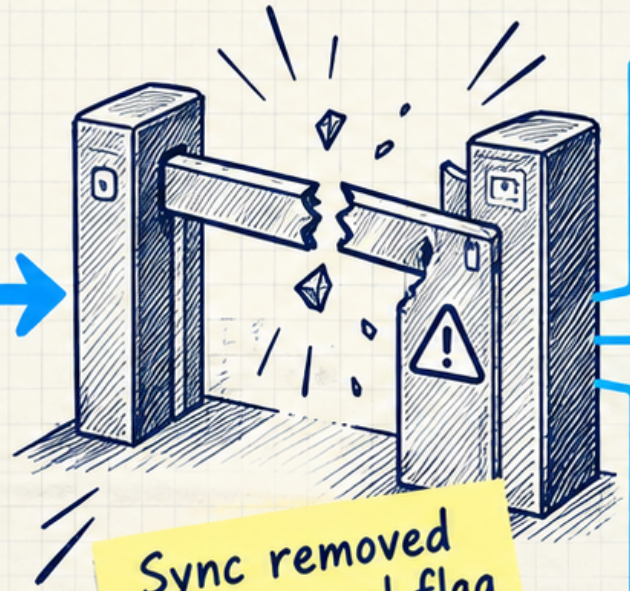

Sync removed the approval flag

RESPONSE: THREE CONTROLS

01 TAXONOMY BUILT
One classification source

02 ADMISSION + CI BUILT
Reject missing gates

03 DAILY AUDIT BUILT
Compare policy with calls

A declared gate needs runtime evidence.

Historical evidence, not a claim about current call counts.

CLOSE THE LOOP, IN ORDER

All five landed, in the order drawn

01 EVALUATION AS A GATE — BUILT • 13 OCT

Post-deploy gate; one agent, eight questions

02 CALLERS AS PRINCIPALS — BUILT • 13 OCT

Verified at the gateway; chat UI still shared

03 DISPATCHER — BUILT • 12 OCT

Alerts → read-only investigation
Next; checks, issues, triage labels

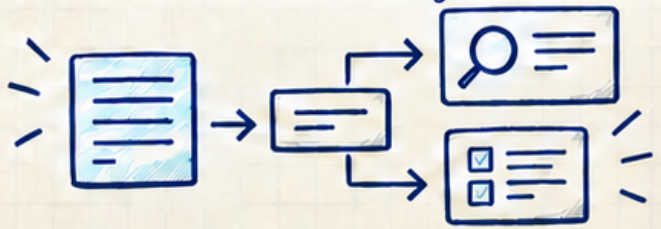

Identity first; all gates held

Evaluation gate

04 IMPLEMENT IN THE PLATFORM — BUILT • 14 OCT

Agent PRs through checks and policy; no sandbox egress

Prior gates required

05 REMEDIATE ON SIGNAL — BUILT • 14 OCT

PR for a low limit; an approved restart for a leak

BUILT REUSE THE FOUNDATION — Secrets • Policy • Docs • Checks • Audit

Next: route checks and issues; a path registry; a retry budget

Which path should earn the next level of autonomy?